


Neuron network with Keras library

minhtt@vietinbank.vn



1. Neuron network library
2. Keras
 - a. Overview
 - b. Model
 - c. Layer
 - d. Optimizer
 - e. Losses
 - f. Callback
 - g. Metric
3. Example
4. Lab
5. Reference

TOP libraries


- Tensor flow
- Pytorch
- MXNet
- ...

Why we need library?

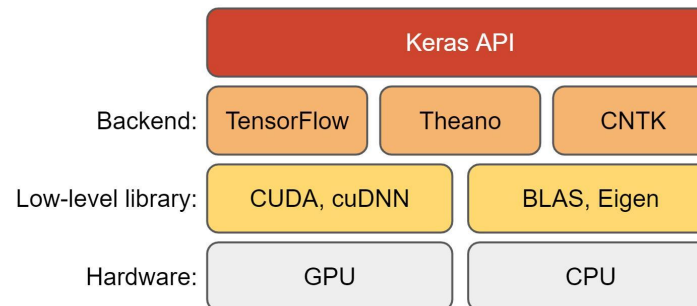
- Easy to build complex neuron network
- Auto gradient
- Utilities to control training and prediction process
- Easy to intergrate with GPU
- Processing unit is tensor



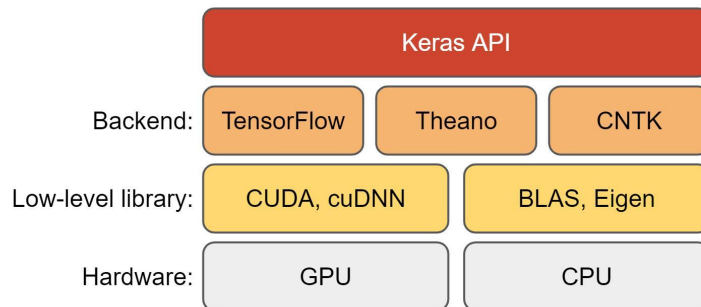
Scalar Vector Matrix Tensor



- Simple
- Flexible
- Powerful



- **Sequential model:** single input, single output
- **Functional API:** support most cases
- **Model subclassing:** out of the box research use cases



A Sequential model is appropriate for a plain stack of layers where each layer has exactly one input tensor and one output tensor.

```
model = keras.Sequential()  
model.add(layers.Dense(2, activation="relu"))  
model.add(layers.Dense(3, activation="relu"))  
model.add(layers.Dense(4))
```

```
model = keras.Sequential(  
    [  
        layers.Dense(2, activation="relu"),  
        layers.Dense(3, activation="relu"),  
        layers.Dense(4),  
    ]  
)
```

- Layer
- A Layer instance is callable, much like a function
- Layers maintain a state, updated when the layer receives data during training, and stored in `layer.weights`
- Types of layers:
 - Core layer
 - Convolution layer
 - Pooling layer
 - Recurrent layer
 - Preprocessing layer
 - Normalization layer
 - Attention layer
 - Reshaping layer
 - Locally connected layer
 - Activation layer

```
from tensorflow.keras import layers

layer = layers.Dense(32, activation='relu')
inputs = tf.random.uniform(shape=(10, 20))
outputs = layer(inputs)
```

- Optimizer
- SGD
- RMSprop
- Adam
- Adadelta
- Adagrad
- Adamax
- NadamFtrl

```
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential()
model.add(layers.Dense(64, kernel_initializer='uniform', input_shape=(10,)))
model.add(layers.Activation('softmax'))

opt = keras.optimizers.Adam(learning_rate=0.01)
model.compile(loss='categorical_crossentropy', optimizer=opt)
```


- Losses

- The purpose of loss functions is to compute the quantity that a model should seek to minimize during training
- Reference in <https://keras.io/api/losses/>

```
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential()
model.add(layers.Dense(64, kernel_initializer='uniform', input_shape=(10,)))
model.add(layers.Activation('softmax'))

loss_fn = keras.losses.SparseCategoricalCrossentropy()
model.compile(loss=loss_fn, optimizer='adam')
```

- Call back
- Write TensorBoard logs after every batch of training to monitor your metrics
- Periodically save your model to disk
- Do early stopping
- Get a view on internal states and statistics of a model during training
- ...

```
my_callbacks = [  
    tf.keras.callbacks.EarlyStopping(patience=2),  
    tf.keras.callbacks.ModelCheckpoint(filepath='model.{epoch:02d}-{val_loss:.2f}.h5'),  
    tf.keras.callbacks.TensorBoard(log_dir='./logs'),  
]  
model.fit(dataset, epochs=10, callbacks=my_callbacks)
```

Predefine metric

```
model.compile(  
    optimizer='adam',  
    loss='mean_squared_error',  
    metrics=[  
        metrics.MeanSquaredError(),  
        metrics.AUC(),  
    ]  
)
```

Custom metric

```
def my_metric_fn(y_true, y_pred):  
    squared_difference = tf.square(y_true - y_pred)  
    return tf.reduce_mean(squared_difference, axis=-1) # Note the `axis=-1`  
  
model.compile(optimizer='adam', loss='mean_squared_error', metrics=[my_metric_fn])
```

- <https://analyticsindiamag.com/top-7-python-neural-network-libraries-for-developers/>
- <https://keras.io/>

The image features a teal background with a bokeh effect of colorful confetti in shades of pink, orange, yellow, and blue. The text "Thank You!" is written in a dark blue, elegant script font, centered horizontally and slightly lower in the frame.

Thank You!